



Yet Another Stupid Simple AI agent.
Local-first, evidence-gated, and ruthlessly token-efficient.

#1

Live Track-1 leaderboard rank

1,292

Fireworks tokens (2nd place:
1,797)

0.947

Judge accuracy (18/19)

13/19

Tasks answered locally, zero
tokens

Go 1.26

2x fine-tuned MiniCPM5-1B

llama.cpp in-container

Granite ModernBERT ONNX


MiniMax-M3 insurance call

4 GB / 2 vCPU fit

Track 1 ranks the fewest Fireworks tokens among submissions that clear the accuracy gate. Local model inference counts as zero.

Local first. Verified always. One insurance call.

Phases peak sequentially, never together: classify, free the classifier, run one model lane at a time.



Two 1B fine-tunes, each owning what it can prove.

Tool lane - MiniCPM5-yassai-v2e3b

- + Maths and logic through a trained run_python contract
- + Code executes in a real subprocess; answers interpolate computed values
- + Training data is parameterised and execution-verified at build time
- + 10/11 unseen maths variants, 6/6 logic with gates

Assist lane - MiniCPM5-yassai-assist-v6

- + Judge-filtered teacher distillation (203/210 kept)
- + Claude-authored adversarial data: sarcasm, idioms, word caps, verified bug-fix pairs
- + Parametric NER/codegen with derived ground truth
- + Held-out behaviour gate before every export - including code functions never trained on

Six fine-tune iterations in one day: train (~2 min on Modal H100), gate, export GGUF, verify against 43 adversarial variants. The pipeline is the moat, not any single checkpoint.

factual

maths

sentiment

summarisation

NER

debugging

logic

code generation

Never trust a free answer. Verify it.

Local answers are free – and worthless unless something deterministic vouches for them.

Grounding

- + Answer numbers must derive from executed stdout
- + Two-vehicle physics bounds, magnitude sanity
- + Generated code must pass the prompt's own worked example
- + Bug fixes must parse, keep the name, and differ from the buggy source

Format and labels

- + Exact bullet/sentence counts, per-bullet word caps
- + Acronym spans are never LOCATIONS; non-dates never DATES
- + Entity recall: years, names, acronyms must appear
- + Sentiment and NER must agree with themselves across two phrasings

Recovery ladder

- + Gate reject -> one gate-guided local retry (reason appended)
- + Still rejected -> folded remote batch
- + Remote gap -> one lean retry, missing ids only
- + Remote dead -> best rejected local answer ships

Measured, not assumed: with gates bypassed the same models score 12/19 – below the accuracy gate. The gates ARE the accuracy. A zero-token mode exists behind a flag for the day a local model clears the bar alone.

6,356 to 1,292 in one day of measured cuts.

Live and verified run history



What actually moved the score

- + Fixed GGUF tokenisation: manual training render + EOG exceptions
- + Both lanes run before batching - no more stranded solo calls
- + Tools only for dedicated code batches (the coin-flip tool loop cost ~2,600)
- + Factual, summaries, bug fixes moved local behind gates
- + Recipes tuned per batch kind; recovery retries stripped lean

rehearsal (4 GB / 2 vCPU emulation): 1,297 tokens
live judge run: 1,292 tokens

The CI dress rehearsal predicted the live score within 5 tokens. Every submission is rehearsed on judge-shaped hardware before it spends a slot.

A reproducible submission, not a notebook demo.

Deterministic CI

- + Native amd64 tests incl. real ONNX integration
- + Image build with versioned GGUF artefacts
- + Dress rehearsal: full task set, 4 GB / 2 vCPU
- + Hard-fails on any empty answer or fallback

Portable image

- + linux/amd64 multi-stage build, immutable SHA tags
- + Two GGUFs + fp32 classifier + llama.cpp b9948
- + Env-only runtime config per the guide
- + Compressed size far under the 10 GB cap

Telemetry loop

- + Every run reports tasks, answers, metrics, traces
- + Cloudflare Worker dashboard, live + rehearsal
- + Diagnosed the timeout roulette and tool-loop spikes
- + Robustness audit: off-distribution wildcard set

#1

1,292 tokens @ 0.947

+505

Token lead over 2nd place

0

Fallbacks, empty answers, DQs

ready

Commit, image, docs, deck